

CS 411: Project 1 Stage 3

For this project we are using Google Cloud Platform (GCP) to host our database. Our database consists of 5 tables provided in Image 1.

```
Database changed
mysql> show tables
-> ;
+-----+
| Tables_in_proj |
+-----+
| combined_phrases |
| country_languages |
| dishes |
| friends |
| monuments |
| users |
+-----+
6 rows in set (0.00 sec)
```

Image 1: Tables that make up the Database

We were successfully able to connect & host the database tables on GCP as showcased in Image 2 below.

```
Welcome to Cloud Shell! Type "help" to get started.
Your Cloud Platform project in this session is set to final-project-cs-411.
Use "gcloud config set project [PROJECT_ID]" to change to a different project.
pranavrajkumar2003@cloudshell:~ (final-project-cs-411)$ gcloud sql connect final-project-cs-411 --user=root
Allowlisting your IP for incoming connection for 5 minutes...done.
Connecting to database with SQL user [root].Enter password:
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 24919
Server version: 8.0.31-google (Google)

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> use proj;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
```

Image 2: A successful connection to the database on GCP

The DDL commands we used for our tables are shown below in Image 3.

```
-- This file contains the English Phrases, their translations, and their
pronunciations to the language of the user's choice.
CREATE TABLE combined_phrases (
    id INTEGER PRIMARY KEY AUTO_INCREMENT,
    english_phrase TEXT NOT NULL,
    translation TEXT NOT NULL,
    pronunciation TEXT NOT NULL,
    language TEXT NOT NULL
);

--The following table is used to store the data for the dishes for each city in the
countries.
CREATE TABLE dishes (
    City VARCHAR(255),
    Country VARCHAR(255) NOT NULL,
    Dish VARCHAR(255) NOT NULL,
    DishID VARCHAR(255) PRIMARY KEY
);

-- The following table is used to store the data for the monuments for each city in
the countries.
CREATE TABLE monuments (
    City VARCHAR(255),
    Country VARCHAR(255) NOT NULL,
    Monument VARCHAR(255) NOT NULL,
    MonumentID VARCHAR(255) PRIMARY KEY
);

-- The following table is used to store the information about the users.
CREATE TABLE Users (
    username VARCHAR(50) NOT NULL PRIMARY KEY,
    points INT NOT NULL,
    first_name VARCHAR(50) NOT NULL,
    last_name VARCHAR(50),
    email VARCHAR(100) NOT NULL,
    Language VARCHAR(50) NOT NULL
);

-- The following table is used to store the information about the friends of the
users.
CREATE TABLE friends (
```

```

friendname1 VARCHAR(50) NOT NULL,
friendname2 VARCHAR(50) NOT NULL,
PRIMARY KEY (friendname1, friendname2),
FOREIGN KEY (friendname1) REFERENCES users(username),
FOREIGN KEY (friendname2) REFERENCES users(username)
);

-- The following table is used to store the information about the countries and their
languages.
CREATE TABLE country_languages (
    Country VARCHAR(50) PRIMARY KEY,
    Language VARCHAR(50) not null
);

```

Image 3: DDL Commands to generate tables

<pre> mysql> select count(*) from combined_phrases; +-----+ count(*) +-----+ 1893 +-----+ 1 row in set (0.08 sec) </pre>	<pre> mysql> select count(*) from dishes; +-----+ count(*) +-----+ 1000 +-----+ 1 row in set (0.00 sec) </pre>
<pre> mysql> select count(*) from monuments; +-----+ count(*) +-----+ 1000 +-----+ 1 row in set (0.00 sec) </pre>	

Image 4: 3 tables of length 1000

The Advanced Queries that were used along with their outputs are provided below in Images 5 - 11.

Query 1:

```
SELECT dish, city, country
FROM dishes
WHERE country in(SELECT L.country
                  FROM users U JOIN country_languages L ON U.Language = L.Language
                  WHERE U.username = 'jstone');
```

Image 5: Advanced Query used to get the Dishes, City & Country based on a language the user is trying to learn

```
mysql> SELECT dish, city, country FROM dishes WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'jstone') Limit 15;
```

dish	city	country
Fondue	San Carlos de Bariloche	Argentina
Trucha Patag�nica	San Carlos de Bariloche	Argentina
Cordero Patag�nico	San Carlos de Bariloche	Argentina
Chocolates Artesanales	San Carlos de Bariloche	Argentina
Goulash con Sp�tztze	San Carlos de Bariloche	Argentina
Ahumados Patag�nicos	San Carlos de Bariloche	Argentina
Empanadas de Cordero	San Carlos de Bariloche	Argentina
Pizza a la Piedra	San Carlos de Bariloche	Argentina
K�schen	San Carlos de Bariloche	Argentina
Helado de Rosa Mosqueta	San Carlos de Bariloche	Argentina
Asado	Buenos Aires	Argentina
Milanesa	Buenos Aires	Argentina
Empanadas	Buenos Aires	Argentina
Provoleta	Buenos Aires	Argentina
Pizza a la Piedra	Buenos Aires	Argentina

```
15 rows in set (0.00 sec)
```

Image 6: Output of Advanced Query 1

Query 2:

```
SELECT Monument, city, country
FROM monuments
WHERE country in(SELECT L.country
                  FROM users U JOIN country_languages L ON U.Language = L.Language WHERE
U.username = 'tlong');
```

Image 7: Advanced Query used to get the Monument, City & Country based on a language the user is trying to learn

```
mysql> SELECT Monument, city, country
-> FROM monuments
-> WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'tlong') limit 15;
+-----+-----+-----+
| Monument | city | country |
+-----+-----+-----+
| Zytglogge (Clock Tower) | Bern | Switzerland |
| Federal Palace of Switzerland | Bern | Switzerland |
| Bear Pit | Bern | Switzerland |
| Bern Cathedral (Münster) | Bern | Switzerland |
| Rosengarten (Rose Garden) | Bern | Switzerland |
| Bundesplatz (Federal Square) | Bern | Switzerland |
| Einstein House | Bern | Switzerland |
| Historical Museum of Bern | Bern | Switzerland |
| Nydegg Bridge | Bern | Switzerland |
| Kornhaus Bridge | Bern | Switzerland |
| Basel Minster | Basel | Switzerland |
| Tinguely Fountain | Basel | Switzerland |
| Rathaus (Town Hall) | Basel | Switzerland |
| Spalentor (Spalen Gate) | Basel | Switzerland |
| Kunstmuseum Basel | Basel | Switzerland |
+-----+-----+-----+
15 rows in set (0.01 sec)
```

Image 8: Output for Advanced Query 2

Query 3:

```
SELECT potential_friend.friendname2 AS suggested_friend,
COUNT(mutual_friends.friendname2) AS mutual_friend_count
FROM friends AS direct_friends JOIN friends AS mutual_friends ON
direct_friends.friendname2 = mutual_friends.friendname1
JOIN friends AS potential_friend ON mutual_friends.friendname2 =
potential_friend.friendname2
WHERE direct_friends.friendname1 = 'aJoshua' AND potential_friend.friendname2 !=
'aJoshua' AND potential_friend.friendname2
NOT IN ( SELECT friendname2 FROM friends WHERE friendname1 = 'aJoshua' )
GROUP BY potential_friend.friendname2 HAVING COUNT(mutual_friends.friendname2) >= 1
ORDER BY mutual_friend_count DESC;
```

Image 9: Get the users who are friends with the friends of the current user and are trying to learn the same language

```
mysql> SELECT potential_friend.friendname2 AS suggested_friend, COUNT(mutual_friends.friendname2) AS mutual_friend_count
-> FROM friends AS direct_friends JOIN friends AS mutual_friends ON direct_friends.friendname2 = mutual_friends.friendname1
-> JOIN friends AS potential_friend ON mutual_friends.friendname2 = potential_friend.friendname2
-> WHERE direct_friends.friendname1 = 'aJoshua' AND potential_friend.friendname2 != 'aJoshua' AND potential_friend.friendname2
-> NOT IN ( SELECT friendname2 FROM friends WHERE friendname1 = 'aJoshua' )
-> GROUP BY potential_friend.friendname2 HAVING COUNT(mutual_friends.friendname2) >= 1
-> ORDER BY mutual_friend_count DESC;
+-----+-----+
| suggested_friend | mutual_friend_count |
+-----+-----+
| pJon | 1 |
| tlong | 1 |
+-----+-----+
2 rows in set (0.05 sec)
```

Image 10: Output for Advanced Query 3

Query 4:

```
SELECT COUNT(DISTINCT m.MonumentID) AS num_monuments, COUNT(DISTINCT d.DishID) AS
num_dishes, m.City AS city
FROM monuments AS m JOIN dishes AS d ON m.City = d.City
GROUP BY m.City;
```

Image 11: Advanced Query to get the count of all the monuments and dishes that have common cities as a fun fact to display to the users.

```
mysql> SELECT COUNT(DISTINCT m.MonumentID) AS num_monuments,
-> COUNT(DISTINCT d.DishID) AS num_dishes,
-> m.City AS city
-> FROM monuments AS m
-> JOIN dishes AS d ON m.City = d.City
-> GROUP BY m.City limit 15;
+-----+-----+-----+
| num_monuments | num_dishes | city      |
+-----+-----+-----+
| 10 | 10 | Åkvara    |
| 10 | 10 | Antwerp   |
| 10 | 10 | Aveiro    |
| 10 | 10 | Barcelona |
| 10 | 10 | Basel     |
| 10 | 10 | Belo Horizonte |
| 10 | 10 | Berlin    |
| 10 | 10 | Bern      |
| 10 | 10 | Bilbao    |
| 10 | 10 | Bordeaux  |
| 10 | 10 | Braga     |
| 10 | 10 | Brasília  |
| 10 | 10 | Bremen    |
| 10 | 10 | Bruges     |
| 10 | 10 | Brussels  |
+-----+-----+-----+
15 rows in set (0.01 sec)
```

Image 12: Output to Advanced Query 4

Indexing Outputs:

Monuments:

Output with no indexing:

```
mysql> explain analyze SELECT dish, city, country
-> FROM dishes
-> WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'jstone');
+-----+-----+-----+
| EXPLAIN |
+-----+-----+-----+
| -> Filter: (dishes.Country = L.Country) (cost=114.10 rows=110) (actual time=0.273..0.864 rows=300 loops=1)
| -> Inner hash join (<hash>(dishes.Country)=<hash>(L.Country)) (cost=114.10 rows=110) (actual time=0.272..0.804 rows=300 loops=1)
| -> Table scan on dishes (cost=11.60 rows=1000) (actual time=0.087..0.422 rows=1000 loops=1)
| -> Hash
| -> Filter: (L.`Language` = 'Spanish') (cost=1.35 rows=1) (actual time=0.122..0.129 rows=3 loops=1)
| -> Table scan on L (cost=1.35 rows=11) (actual time=0.063..0.069 rows=11 loops=1)
|
+-----+-----+-----+
1 row in set (0.02 sec)
```

Index on country_languages.Language:

```
mysql> CREATE
-> INDEX id_country_languages_language ON country_languages (Language) ;
Query OK, 0 rows affected (0.23 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> explain analyze SELECT dish, city, country FROM dishes WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'jstone');
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Filter: (dishes.Country = L.Country) (cost=303.34 rows=300) (actual time=0.104..0.635 rows=300 loops=1)
|   -> Inner hash join (hashb(dishes.Country)=hashb(L.Country)) (cost=303.34 rows=300) (actual time=0.102..0.572 rows=300 loops=1)
|     -> Table scan on dishes (cost=4.25 rows=1000) (actual time=0.031..0.368 rows=1000 loops=1)
|       -> Hash
|         -> Covering index lookup on L using id_country_languages_language (Language='Spanish') (cost=0.57 rows=3) (actual time=0.019..0.022 rows=3 loops=1)
|
+-----+
1 row in set (0.01 sec)
```

We wanted to add an index on Language in country_languages to help the join process but as you can see this increased the cost even though it reduced the time for the filtering. This was not the best move and we removed this indexing.

Index on dishes.Country:

```
mysql> CREATE INDEX idx_dishes_country ON dishes(Country);
Query OK, 0 rows affected (0.16 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> explain analyze SELECT dish, city, country FROM dishes WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'jstone');
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Nested loop inner join (cost=21.43 rows=110) (actual time=0.887..12.555 rows=300 loops=1)
|   -> Filter: (L.Language = 'Spanish') (cost=1.35 rows=1) (actual time=0.015..0.031 rows=3 loops=1)
|     -> Table scan on L (cost=1.35 rows=11) (actual time=0.013..0.023 rows=11 loops=1)
|       -> Index lookup on dishes using idx_dishes_country (Country=L.Country), with index condition: (dishes.Country = L.Country) (cost=17.34 rows=100) (actual time=4.012..4.166 rows=100 loops=3)
|
+-----+
1 row in set (0.02 sec)
```

Instead we decided to add indexing on Country in dishes to help the WHERE clause in the main query filter rows from dishes by Country. As you can see this has helped in reducing the cost and we decided to keep this indexing.

Monuments:

Output with no indexing:

```
mysql> explain analyze SELECT Monument, city, country
-> FROM monuments
-> WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'tlong')
-> ;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Filter: (monuments.Country = L.Country) (cost=114.10 rows=110) (actual time=0.311..0.682 rows=200 loops=1)
-> Inner hash join (<hash>(monuments.Country)=<hash>(L.Country)) (cost=114.10 rows=110) (actual time=0.311..0.636 rows=200 loops=1)
-> Table scan on monuments (cost=11.60 rows=1000) (actual time=0.041..0.403 rows=1000 loops=1)
-> Hash
-> Filter: (L.Language = 'German') (cost=1.35 rows=1) (actual time=0.055..0.058 rows=2 loops=1)
-> Table scan on L (cost=1.35 rows=11) (actual time=0.027..0.030 rows=11 loops=1)
|
+-----+
1 row in set (0.01 sec)
```

Index on country_languages.Language:

```
mysql> CREATE INDEX idx_country_languages_language ON country_languages (Language);
Query OK, 0 rows affected (0.07 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> explain analyze SELECT Monument, city, country FROM monuments WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'tlong');
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Filter: (monuments.Country = L.Country) (cost=203.22 rows=200) (actual time=0.369..0.732 rows=200 loops=1)
-> Inner hash join (<hash>(monuments.Country)=<hash>(L.Country)) (cost=203.22 rows=200) (actual time=0.367..0.675 rows=200 loops=1)
-> Table scan on monuments (cost=6.38 rows=1000) (actual time=0.035..0.400 rows=1000 loops=1)
-> Hash
-> Covering index lookup on L using idx_country_languages_language (Language='German') (cost=0.46 rows=2) (actual time=0.013..0.017 rows=2 loops=1)
|
+-----+
1 row in set (0.00 sec)
```

We wanted to add an index on Language in country_languages to help the join process but as you can see this increased the cost even though it reduced the time for the filtering. This was not the best move and we removed this indexing.

Index on monuments.Country:

```
mysql> CREATE INDEX idx_monuments_country ON monuments (Country);
Query OK, 0 rows affected (0.13 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> explain analyze SELECT Monument, city, country FROM monuments WHERE country in(SELECT L.country FROM users U JOIN country_languages L ON U.Language = L.Language WHERE U.username = 'tlong');
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Nested loop inner join (cost=21.43 rows=110) (actual time=0.164..10.083 rows=200 loops=1)
-> Filter: (L.Language = 'German') (cost=1.35 rows=1) (actual time=0.015..0.024 rows=2 loops=1)
-> Table scan on L (cost=1.35 rows=11) (actual time=0.011..0.020 rows=11 loops=1)
-> Index lookup on monuments using idx_monuments_country (Country=L.Country), with index condition: (monuments.Country = L.Country) (cost=17.34 rows=100) (actual time=4.808..5.021 rows=100 loops=2)
|
+-----+
1 row in set (0.02 sec)
```

Instead we decided to add indexing on Country in monuments to help the WHERE clause in the main query filter rows from dishes by Country. As you can see this has helped in reducing the cost and we decided to keep this indexing.

Friends:

Output with no indexing:

[illegible]

Output with Index on friends (friendname1, friendname2):

[illegible]

We added this indexing as we thought this would help the joins perform faster by covering both friendname1 and friendname2 columns and as you can see it has reduced the time but not the cost.

Output with Index on friends (friendname2) and friends (friendname1):

```
mysql> CREATE INDEX idx_friendname1 ON friends(friendname1);
Query OK, 9 rows affected (0.04 sec)
Records: 9 Duplicates: 0 Warnings: 0

mysql> explain analyze SELECT potential_friend.friendname2 AS suggested_friend, COUNT(mutual_friends.friendname2) AS mutual_friend_count FROM friends AS direct_friends JOIN friends AS mutual_friends ON direct_friends.friendname2 = mutual_friends.friendname1 JOIN friends AS potential_friends ON mutual_friends.friendname1 = potential_friend.friendname2 WHERE direct_friends.friendname1 = 'aJoshua' AND potential_friend.friendname2 != 'aJoshua' AND potential_friend.friendname2 NOT IN ( SELECT friendname2 FROM friends WHERE friendname1 = 'aJoshua' ) GROUP BY potential_friend.friendname2 HAVING COUNT(mutual_friends.friendname2) >= 1 ORDER BY mutual_friend_count DESC;
+----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | Extra |
+----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | friends | | index | PRIMARY | PRIMARY | 1 | | 9 | Using index |
+----+-----+-----+-----+-----+-----+
EXPLAIN
+----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | Extra |
+----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | friends | | index | PRIMARY | PRIMARY | 1 | | 9 | Using index |
+----+-----+-----+-----+-----+-----+

--> Sort: mutual_friend_count DESC (actual time=0.073..0.073 rows=2 loops=1)
--> Filter: (count(mutual_friends.friendname2) >= 1) (actual time=0.062..0.062 rows=2 loops=1)
--> Table scan on <tempstorage> (actual time=0.039..0.039 rows=2 loops=1)
--> Aggregate using temporary table (actual time=0.057..0.057 rows=2 loops=1)
--> Nested loop antijoin (cost=1.97 rows=2) (actual time=0.032..0.039 rows=2 loops=1)
--> > Nested loop inner join (cost=0.32 rows=2) (actual time=0.028..0.034 rows=2 loops=1)
--> > > Nested loop inner join (cost=0.76 rows=2) (actual time=0.025..0.028 rows=2 loops=1)
--> > > > Covering index lookup on direct_friends using PRIMARY (friendname1='aJoshua') (cost=0.35 rows=1) (actual time=0.016..0.016 rows=1 loops=1)
--> > > > > Filter: (mutual_friends.friendname2 <> 'aJoshua') (cost=0.41 rows=2) (actual time=0.008..0.009 rows=2 loops=1)
--> > > > > > Covering index lookup on mutual_friends using idx_friendname1 (friendname1=direct_friends.friendname2) (cost=0.41 rows=2) (actual time=0.006..0.008 rows=1 loops=1)
--> > > > > > > Covering index lookup on potential_friends using idx_friendname2 (friendname2=mutual_friends.friendname2) (cost=0.38 rows=2) (actual time=0.002..0.004 rows=1 loops=2)
--> > > > > > > > Single row covering index lookup on friends using PRIMARY (friendname1='aJoshua', friendname2=mutual_friends.friendname2) (cost=0.21 rows=1) (actual time=0.002..0.002 rows=0 loops=2)

1 row in set (0.00 sec)
```

We added an index on friends(friendname2) to speed up the anti-join subquery, especially for potential_friend retrievals and we can see this is happening from the images as the time is reducing. Similarly index on friends(friendname1) enables efficient scans on friendname1 for quick filtering of relevant users.

Count Information:

Output with no indexing:

```
mysql> EXPLAIN ANALYZE
-> SELECT COUNT(DISTINCT m.MonumentID) AS num_monuments,
->        COUNT(DISTINCT d.DishID) AS num_dishes,
->        m.City AS city
-> FROM monuments AS m
-> JOIN dishes AS d ON m.City = d.City
-> GROUP BY m.City;
+-----+
| EXPLAIN
+-----+
|
+-----+
| -> Group aggregate: count(distinct monuments.MonumentID), count(distinct dishes.DishID) (actual time=9.724..19.263 rows=99 loops=1)
| -> Sort: m.City (actual time=9.619..10.238 rows=10200 loops=1)
|   -> Stream results (cost=100126.95 rows=100000) (actual time=0.648..5.510 rows=10200 loops=1)
|     -> Filter: (d.City = m.City) (cost=100126.95 rows=100000) (actual time=0.643..3.183 rows=10200 loops=1)
|       -> Inner hash join (<hash>(d.City)=<hash>(m.City)) (cost=100126.95 rows=100000) (actual time=0.640..1.522 rows=10200 loops=1)
|         -> Table scan on d (cost=0.03 rows=1000) (actual time=0.023..0.270 rows=1000 loops=1)
|         -> Hash
|           -> Table scan on m (cost=102.75 rows=1000) (actual time=0.039..0.405 rows=1000 loops=1)
|
+-----+
1 row in set (0.03 sec)
```

Output with index on monument.City and dish.City:

```
mysql> CREATE INDEX idx_monuments_city ON monuments(City);
Query OK, 0 rows affected (0.20 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> CREATE INDEX idx_dishes_city ON dishes(City);
Query OK, 0 rows affected (0.08 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
| -> Group aggregate: count(distinct m.MonumentID), count(distinct d.DishID) (cost=3885.58 rows=10101) (actual time=2.155..15.944 rows=99 loops=1)
| -> Nested loop inner join (cost=2875.48 rows=10101) (actual time=1.942..7.305 rows=10200 loops=1)
|   -> Filter: (m.City is not null) (cost=102.75 rows=1000) (actual time=1.826..2.084 rows=1000 loops=1)
|     -> Covering index scan on m using idx_monuments_city (cost=102.75 rows=1000) (actual time=0.825..1.016 rows=1000 loops=1)
|       -> Covering index lookup on d using idx_dishes_city (City=m.City) (cost=1.76 rows=10) (actual time=0.003..0.004 rows=10 loops=1000)
```

We added indexes on monument.City and dish.City as the query joins monuments and dishes on City and indexing City in both tables should optimize the join process. As we can see the cost has significantly reduced and we decided to keep this indexing.

Output with index on monument.MonumentsID and dish.DishID:

```
mysql> CREATE INDEX idx_monuments monumentid ON monuments(MonumentID);
Query OK, 0 rows affected (0.13 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> CREATE INDEX idx_dishes dishid ON dishes(DishID);
Query OK, 0 rows affected (0.09 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql> EXPLAIN ANALYZE SELECT COUNT(DISTINCT m.MonumentID) AS num_monuments,
                                COUNT(DISTINCT d.DishID) AS num_dishes,
                                m.City AS city FROM monuments AS m JOIN dishes AS d ON m.City = d.City GROUP BY m.City;
+-----+
| EXPLAIN |
+-----+
+-----+
+-----+
+-----+
+-----+
-> Group aggregate: count(distinct m.MonumentID), count(distinct d.DishID) (cost=3885.58 rows=10101) (actual time=0.375..17.632 rows=99 loops=1)
-> Nested loop inner join (cost=273.48 rows=10101) (actual time=0.000..6.482 rows=1000 loops=1)
    -> Filter: (m.City is not null) (cost=102.75 rows=1000) (actual time=0.049..0.471 rows=1000 loops=1)
        -> Covering index scan on m using idx_monuments_city (cost=102.75 rows=1000) (actual time=0.047..0.389 rows=1000 loops=1)
            -> Covering index lookup on d using idx_dishes_city (City=m.City) (cost=1.74 rows=10) (actual time=0.003..0.005 rows=10 loops=1000)
```

The count operations can be sped up by indexing MonumentID and DishID in the monuments and dishes table respectively and looking at the output we can see the time has reduced and the cost has not changed.

Updates to Database Design Based on Stage 2 Feedback

Based on the feedback, we combined the tables for all languages instead of having a separate table for each language. We also accounted for the self-loop for the friends entity but have kept the table as we have a many to many relationship. We added new entities to meet the 5 entity requirements by adding the new dishes, monuments & countries entities that provide information about the most popular languages, food dishes & monuments in countries. With this addition, we are not only proposing a language-learning application but an overall tourist experience where we will provide users with the language of the city, popular food items and the famous monuments of the city.